RAG ASSIGNMENT REPORT: BASELINE VS. RE-RANKING

Student: Raj K. Jain | Deliverable: Technical Report | Course: LMA Systems

1. ARCHITECTURE COMPARISON: BI-ENCODER VS. CROSS-ENCODER

In modern Retrieval-Augmented Generation (RAG) systems, two fundamentally distinct neural architectures are used to score candidate passages for an incoming query:

A. Bi-Encoder Architecture (Two-Tower / Independent Dense Vectors):

? Operates with two independent encoder towers (or a shared tower evaluated separately). Query q and Chunk c are passed independently.

? Produces dense vector representations: u = E(q) and v = E(c), where u, v in R^384. Semantic similarity is computed via cosine similarity: cos(u, v).

? Critical Advantage: Vector representations of document chunks can be computed offline once and stored in Approximate Nearest Neighbor (ANN) vector indices (FAISS, Milvus, HNSW). Query search scales in O(log N) sub-milliseconds.

? Major Limitation (Representation Bottleneck): Zero token-level cross-attention occurs between query and chunk. Squeezing an entire 300-word passage into a 384-dimensional vector causes severe loss of fine-grained qualifiers, negations, and exceptions.

B. Cross-Encoder Architecture (Joint All-to-All Self-Attention):

? Concatenates query and chunk into a single sequence: [CLS] query [SEP] chunk [SEP], passing it jointly through all transformer layers.

? Every token in the query attends directly to every token in the passage across all self-attention heads at every layer.

? Critical Advantage (SOTA Precision): Deep cross-attention detects nuanced linguistic dependencies, modifier constraints ('only if', 'prohibited'), and negations.
? Major Limitation: Joint representation cannot be precomputed offline. Scoring N chunks requires N full forward passes through the transformer, making full-corpus search computationally prohibitive.

C. The Two-Stage Paradigm (Best of Both Worlds):

Production RAG architectures resolve this trade-off using a Two-Stage Pipeline: Stage 1 utilizes a Bi-Encoder to rapidly retrieve the Top 10 candidates from millions of chunks in ~5 ms. Stage 2 executes Cross-Encoder re-ranking across only those 10 candidate pairs in ~18 ms, selecting the refined Top 3 for LLM generation.

2. STRESS TESTING & FAILURE MODE ANALYSIS

To evaluate the practical benefit of Cross-Encoder re-ranking, we tested queries targeting specific failure modes where baseline dense retrieval stumbles:

Failure Mode 1: Exceptions & Corner Cases (Buried at Rank #4)

? Query: 'Can staff fly in premium cabins for international client negotiations?'
? Target Truth: Section 2.4 ('Long-Haul Exception Clause') permits Business Class for international flights >6 hours if presenting to clients within 12 hours, with VP/CFO approval.
? Baseline Dense Failure: Dense retrieval retrieved the general rule in Section 2.1 ('All flights must be Standard Economy') at Rank #1, and buried Section 2.4 at Rank #4 (outside the Top 3!).

? Baseline LLM Answer 1: Inaccurately stated: 'No, all travel must be in Standard Economy. No exceptions exist for client meetings.' (Severe policy misstatement).
? Cross-Encoder Resolution: Promoted Section 2.4 from Rank #4 -> Rank #1 (+3 ranks). Re-ranked Answer 2 accurately stated the exact long-haul client negotiation exception!

Failure Mode 2: Negation & Constraints (Buried at Rank #4)

? Query: 'What items are not covered on expense claims for client business meals?'
? Target Truth: Section 5.3 strictly limits alcohol to 1 drink per diner, bans solo alcohol 100%, and prohibits nightclub bottle service, adult venues, and tobacco.
? Baseline Dense Failure: Bi-Encoder matched positive client dinner text (Section 5.1) at Rank #1 and buried Section 5.3 at Rank #4.

? Baseline LLM Answer 1: Mentioned allowed meals and tips, completely omitting prohibited items and alcohol caps.

? Cross-Encoder Resolution: Detected negation markers and promoted Section 5.3 from Rank #4 -> Rank #2 (+2 ranks), giving the comprehensive list of prohibited items.

Failure Mode Summary Table:

? TEST_01 (Exceptions & Corner Cases): Baseline Rank #4 -> Re-Ranked Rank #1 (Shift: +3 ranks)
? TEST_02 (Negation & Constraints): Baseline Rank #4 -> Re-Ranked Rank #2 (Shift: +2 ranks)
? TEST_03 (Keyword Distractor): Baseline Rank #2 -> Re-Ranked Rank #2 (Preserved in Top 2)
? TEST_04 (Buried Penalty Clause): Baseline Rank #1 -> Re-Ranked Rank #1 (Preserved at Rank #1)

LATENCY PROFILING & SCALING ANALYSIS

3. LATENCY PROFILING BENCHMARK (10 WARMED ITERATIONS)

Runtimes recorded on NVIDIA RTX GPU:

Test Case	P1 Baseline	Cross-Encoder (10 pairs)	P2 Total	Overhead Ratio
TEST_01 (Except.)	9.80 ms	17.99 ms	27.79 ms	2.83x
TEST_02 (Negat.)	9.91 ms	21.02 ms	30.93 ms	3.12x
TEST_03 (Distr.)	8.81 ms	19.03 ms	27.84 ms	3.16x
TEST_04 (Penalty)	9.53 ms	19.42 ms	28.95 ms	3.04x
AVERAGE	9.51 ms	19.37 ms	28.88 ms	3.04x

Key Takeaway: Re-ranking adds ~19.4 ms overhead (~3.0x latency factor). In production, an end-to-end retrieval latency of ~29 ms is well within human real-time perception limits (<100 ms) while completely eliminating critical compliance hallucinations.

4. SCALING QUESTION: WHY NOT RUN A CROSS-ENCODER ACROSS THE ENTIRE CORPUS?

A fundamental question in IR system design is why we cannot simply score all document chunks using a Cross-Encoder from the very start:

1. Mathematical Explosion of Compute Complexity:
   ? A Cross-Encoder requires evaluating transformer self-attention over every (query, chunk) pair.

   ? If an enterprise corpus contains N = 1,000,000 chunks, scoring a single query requires 1,000,000 forward passes.

   ? At 2 ms per inference on an enterprise GPU: 1,000,000 x 2 ms = 2,000 seconds ? 33.3 MINUTES per query!

2. Impossibility of Offline Vector Indexing:
   ? In a Bi-Encoder, chunk vectors are independent of the query and precomputed once offline, allowing sublinear ANN search in 5 ms via FAISS/HNSW.

   ? In a Cross-Encoder, chunk representation depends on the query; no offline vector database can ever be pre-built.

3. Conclusion:
   ? The Two-Stage architecture is mathematically necessary: Stage 1 reduces 1,000,000 chunks to 10 in 5 ms, and Stage 2 precisely ranks those 10 chunks in 15 ms, achieving state-of-the-art precision in only 20 ms total!